Ganesh Atmakuri

+15134625201 | atmakuriganesh1234@gmail.com | linkedin.com/in/atmakuriganesh | github.com/atmakuriganesh

Technical Skills

Languages and Database: Python, SQL (PostgreSQL, MySQL), R, Java
Frameworks and libraries: Pandas, Matplotlib, PyTorch, TensorFlow, Sklearn, MLflow, FastAPI, Airflow, Langchain
Tools and Platforms: Power BI, Linux, Git, Docker, Snowflake, AWS Bedrock, Google Vertex AI, Azure Open AI
Development: Data visualization, ETL, Statistical Analysis, ML, DL, NLP, GenAI, Agentic AI (LangGraph, CrewAI)
Certifications: Microsoft Certified: Azure Data Scientist Associate (Dp-100)

Experience

Data Scientist Aug 2023 – Dec 2024
University of Cincinnati *Cincinnati, OH*

- Performed comprehensive exploratory data analysis (EDA) and implemented machine learning models on departmental research data. This effort significantly improved data-driven decision-making, leading to a 15% increase in the efficiency of funding allocation and overall research project outcomes
- Designed and implemented automated data processing pipelines using Python (Pandas, NumPy), incorporating A/B testing that evaluated different feature sets, resulting in 92% reduction in manual review time and 23% improvement in model accuracy
- Developed interactive dashboards using PowerBI to track and report key performance metrics for sponsored research activities, facilitating real-time insights and enabling a 10% reduction in reporting time for stakeholders

Senior Analyst Jan 2022 – July 2023
Capgemini *Benguluru, India*

- Developed and trained ensemble ML models (XGBoost, LightGBM) using scikit-learn and performed Bayesian hyperparameter optimization with Optuna, improving model accuracy by 35% and reducing false positives by 28% across prediction tasks handling large amount of structured data
- Architected end-to-end MLOps pipeline using MLflow for experiment tracking and model registry, implementing automated CI/CD with GitHub Actions and Docker, resulting in 40% faster model deployment cycles and reproducible training workflows across teams
- Designed scalable ML infrastructure on AWS (Sagemaker, Lambda, ECR) with automated model retraining pipelines, A/B testing frameworks, and real-time performance monitoring using CloudWatch, achieving 99.9% system reliability and 45% reduction in inference latency
- Orchestrated distributed model training on AWS for parallel hyperparameter optimization across 50+ experiments, implementing cross-validation strategies and feature selection techniques that improved model F1-score by 32% while reducing training time by 60%

Data Analyst Intern July 2021 – Dec 2021
Swecha AP *Guntur, India*

- Engineered data models using SQL and Python to analyze user interaction patterns across multiple databases, and cohort analysis queries that improved information retention by 15% and increased user engagement by 10%
- Built interactive BI dashboards using PowerBI and Python (Plotly) to visualize simulation effectiveness across 50K+ user sessions, implementing automated anomaly detection and cohort analysis that improved stakeholder satisfaction by 20% and optimized resource allocation through data-driven insights

Projects

Bank Term Deposit Subscription Prediction | *ML, Ensemble Techniques, Exploratory Data Analysis (EDA)*

- Engineered hybrid ML/DL solution combining ensemble methods with deep neural networks (LSTM, Transformer) to predict client term deposit subscriptions, achieving 90%+ accuracy through feature engineering and sequential pattern analysis of banking transactions, resulting in 35% improved targeting efficiency
- Conducted comprehensive exploratory data analysis on a dataset with over 41,000 records, identifying key features that enhanced model performance by 15% compared to traditional classifiers

Student Performance Prediction System | *Python, Scikit-learn, AWS, Azure, Flask*

- Engineered an end-to-end ML pipeline with GridSearchCV for hyperparameter tuning achieving 88% prediction accuracy for student math scores by implementing and evaluating 7+ ML algorithms
- Built and deployed RESTful prediction API using Flask and Docker on both AWS Beanstalk and Azure Container Services, achieving 99.9% uptime and 25ms average response time for 100+ concurrent requests

Education

University of Cincinnati Aug 2023 – Dec 2024
Master of Science in Information Technology, Data science concentration *Cincinnati, OH*